



DevOps external course

Linux administration with Bash. Lektion 1

Lecture 7.1

Module 7 **Linux Administration + Bash**

Serge Prykhodchenko



Agenda

- Iptables – cont.
- Shell
- Scripting basics
- Q&A

iptables

Display a list of rules with line numbers.

iptables -n -L -v --line-numbers

Display the INPUT or OUTPUT chains of rules.

iptables -L INPUT -n -v

iptables -L OUTPUT -n -v --line-numbers

iptables

Stop, start, restart the firewall.

By the forces of the system itself:

```
# service ufw stop
```

```
# service ufw start
```

You can also use the iptables commands to stop the firewall and remove all rules:

```
# iptables -F
```

```
# iptables -X
```

```
# iptables -t nat -F
```

```
# iptables -t nat -X
```

```
# iptables -t mangle -F
```

```
# iptables -t mangle -X
```

```
# iptables -P INPUT ACCEPT
```

```
# iptables -P OUTPUT ACCEPT
```

```
# iptables -P FORWARD ACCEPT
```

Where:

-F: Flush all rules.

-X: Remove the chain.

-t table_name: Select a table (nat or mangle) and remove all rules.

-P: Select default actions (such as DROP, REJECT, or ACCEPT).

iptables

Remove firewall rules.

To display the line number with existing rules:

```
# iptables -L INPUT -n --line-numbers
```

```
# iptables -L OUTPUT -n --line-numbers
```

```
# iptables -L OUTPUT -n --line-numbers | less
```

```
# iptables -L OUTPUT -n --line-numbers | grep 202.54.1.1
```

Let's get a list of IP addresses.

Just look at the number on the left and delete the corresponding line. For example, for number 3:

```
# iptables -D INPUT 3
```

Or find the source IP address (202.54.1.1) and remove it from the rule:

```
# iptables -D INPUT -s 202.54.1.1 -j DROP
```

Where:

-D: Remove one or more rules from the chain.

iptables

Add a rule to the firewall.

To add one or more rules to a chain, we first display a list using line numbers:

```
# iptables -L INPUT -n --line-numbers
```

Approximate output:

Chain INPUT (policy DROP)

<i>num</i>	<i>target</i>	<i>prot</i>	<i>opt</i>	<i>source</i>	<i>destination</i>	
1	DROP	all	--	202.54.1.1	0.0.0.0/0	
2	ACCEPT	all	--	0.0.0.0/0	0.0.0.0/0	state NEW,ESTABLISHED

To insert a rule between 1 and 2 lines:

```
# iptables -I INPUT 2 -s 202.54.1.2 -j DROP
```

iptables

Let's check if the rule has been updated:

```
# iptables -L INPUT -n --line-numbers
```

The output will look like this:

Chain INPUT (policy DROP)

<i>num</i>	<i>target</i>	<i>prot</i>	<i>opt</i>	<i>source</i>	<i>destination</i>	
1	DROP	all	--	202.54.1.1	0.0.0.0/0	
2	DROP	all	--	202.54.1.2	0.0.0.0/0	
3	ACCEPT	all	--	0.0.0.0/0	0.0.0.0/0	<i>state NEW,ESTABLISHED</i>

iptables

Let's save the firewall rules.

Via iptables-save:

```
# iptables-save > /etc/iptables.rules
```

Restoring the rules.

Via iptables-restore

```
# iptables-restore < /etc/iptables.rules
```

set the default policies.

To drop all traffic:

```
# iptables -P INPUT DROP
```

```
# iptables -P OUTPUT DROP
```

```
# iptables -P FORWARD DROP
```

```
# iptables -L -v -n
```

After the above commands, no packet will leave this host.

```
# ping google.com
```


iptables

Block only incoming connections.

To drop all incoming packets not initiated by you, but allow outgoing traffic:

```
# iptables -P INPUT DROP
```

```
# iptables -P FORWARD DROP
```

```
# iptables -P OUTPUT ACCEPT
```

```
# iptables -A INPUT -m state --state NEW,ESTABLISHED -j ACCEPT
```

```
# iptables -L -v -n
```

Outgoing packets and those that were remembered within the established sessions are allowed.

```
# ping google.com
```

iptables

Blocking a specific IP address.

To block the address of the cracker 1.2.3.4:

```
# iptables -A INPUT -s 1.2.3.4 -j DROP
```

```
# iptables -A INPUT -s 192.168.0.0/24 -j DROP
```

Block incoming port requests.

To block all incoming requests on port 80:

```
# iptables -A INPUT -p tcp --dport 80 -j DROP
```

```
# iptables -A INPUT -i eth1 -p tcp --dport 80 -j DROP
```

To block a request for port 80 from address 1.2.3.4:

```
# iptables -A INPUT -p tcp -s 1.2.3.4 --dport 80 -j DROP
```

```
# iptables -A INPUT -i eth1 -p tcp -s 192.168.1.0/24 --dport 80 -j DROP
```

iptables

Block requests to the outgoing IP address.

To block a specific domain, find out its address:

```
# host -t a facebook.com
```

So, facebook.com has address 69.171.228.40

Let's find CIDR for 69.171.228.40:

```
# whois 69.171.228.40 | grep CIDR
```

output:

```
CIDR: 69.171.224.0/19
```

Block access to 69.171.224.0/19:

```
# iptables -A OUTPUT -p tcp -d 69.171.224.0/19 -j DROP
```

You can also use a domain to block:

```
# iptables -A OUTPUT -p tcp -d www.facebook.com -j DROP
```

```
# iptables -A OUTPUT -p tcp -d facebook.com -j DROP
```

iptables

Record event and reset.

To log the movement of packets before dropping, add a rule:

```
# iptables -A INPUT -i eth1 -s 10.0.0.0/8 -j LOG --log-prefix "IP_SPOOF A: "
```

```
# iptables -A INPUT -i eth1 -s 10.0.0.0/8 -j DROP
```

Let's check the log (by default / var / log / messages):

```
# tail -f /var/log/messages
```

```
# grep -i --color 'IP SPOOF' /var/log/messages
```

Record the event and reset (with a limit on the number of records).

In order not to fill the section with a bloated log, we will limit the number of entries using -m.

For example, to write a maximum of 7 lines every 5 minutes:

```
# iptables -A INPUT -i eth1 -s 10.0.0.0/8 -m limit --limit 5/m --limit-burst 7 -j LOG --log-prefix "IP_SPOOF A: "
```

```
# iptables -A INPUT -i eth1 -s 10.0.0.0/8 -j DROP
```

iptables

Drop or allow traffic from specific MAC addresses.

```
# iptables -A INPUT -m mac --mac-source 00:0F:EA:91:04:08 -j DROP
```

```
## * allow only for TCP port # 8080 from mac address 00: 0F: EA: 91: 04: 07 * ##
```

```
# iptables -A INPUT -p tcp --destination-port 22 -m mac --mac-source 00:0F:EA:91:04:07 -j ACCEPT
```

Allow or deny ICMP Ping requests.

To disable ping:

```
# iptables -A INPUT -p icmp --icmp-type echo-request -j DROP
```

```
# iptables -A INPUT -i eth1 -p icmp --icmp-type echo-request -j DROP
```

Allow for specific networks / hosts:

```
# iptables -A INPUT -s 192.168.1.0/24 -p icmp --icmp-type echo-request -j ACCEPT
```

Allow only part of ICMP requests:

```
### ** assumes default inbound policies are set to DROP ** ###
```

```
# iptables -A INPUT -p icmp --icmp-type echo-reply -j ACCEPT
```

```
# iptables -A INPUT -p icmp --icmp-type destination-unreachable -j ACCEPT
```

```
# iptables -A INPUT -p icmp --icmp-type time-exceeded -j ACCEPT
```

```
## ** let us reply to the request ** ##
```

```
# iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT
```

iptables

Open port range.

```
# iptables -A INPUT -m state --state NEW -m tcp -p tcp --dport 7000:7010 -j ACCEPT
```

Open address range.

```
## allow connection to port 80 (Apache) if the address is in the range from 192.168.1.100 to  
192.168.1.200 ##
```

```
# iptables -A INPUT -p tcp --destination-port 80 -m iprange --src-range 192.168.1.100-192.168.1.200 -j  
ACCEPT
```

```
## example for nat ##
```

```
# iptables -t nat -A POSTROUTING -j SNAT --to-source 192.168.1.20-192.168.1.25
```

iptables

Close or open standard ports.

Replace ACCEPT with DROP to block the port.

ssh tcp port 22

iptables -A INPUT -m state --state NEW -m tcp -p tcp --dport 22 -j ACCEPT

iptables -A INPUT -s 192.168.1.0/24 -m state --state NEW -p tcp --dport 22 -j ACCEPT

cups (printing service) udp/tcp port 631 для локальной сети

iptables -A INPUT -s 192.168.1.0/24 -p udp -m udp --dport 631 -j ACCEPT

iptables -A INPUT -s 192.168.1.0/24 -p tcp -m tcp --dport 631 -j ACCEPT

time sync via NTP для локальной сети (udp port 123)

iptables -A INPUT -s 192.168.1.0/24 -m state --state NEW -p udp --dport 123 -j ACCEPT

tcp port 25 (smtp)

iptables -A INPUT -m state --state NEW -p tcp --dport 25 -j ACCEPT

dns server ports

iptables -A INPUT -m state --state NEW -p udp --dport 53 -j ACCEPT

iptables -A INPUT -m state --state NEW -p tcp --dport 53 -j ACCEPT

iptables

Limit the number of parallel connections to the server for one address.

The connlimit module is used for restrictions. To allow only 3 ssh connections per client:

```
# iptables -A INPUT -p tcp --syn --dport 22 -m connlimit --connlimit-above 3 -j REJECT
```

Set the number of HTTP requests to 20:

```
# iptables -p tcp --syn --dport 80 -m connlimit --connlimit-above 20 --connlimit-mask 24 -j DROP
```

Where are:

--connlimit-above 3: Specifies that the rule only takes effect if the number of connections exceeds 3.

--connlimit-mask 24: Specifies the netmask.

iptables

The standard set of rules for a web server on port 80 and SSH:

Default policies

iptables -P INPUT DROP

iptables -P FORWARD DROP

iptables -P OUTPUT ACCEPT

Allow incoming connections from localhost

iptables -A INPUT -i lo -j ACCEPT

Allow established incoming connections

iptables -A INPUT -m state --state ESTABLISHED,RELATED -j ACCEPT

Allow HTTP traffic

iptables -A INPUT -p TCP --dport 80 -j ACCEPT

Allow SSH traffic

iptables -A INPUT -p TCP --dport 22 -j ACCEPT

<https://wiki.dieg.info/iptables>

nftables

nftables is a subsystem of the Linux kernel providing filtering and classification of network packets/datagrams/frames. It has been available since Linux kernel 3.13 released on 19 January 2014.

nftables replaces the legacy iptables portions of Netfilter. Among the advantages of nftables over iptables is less code duplication and easier extension to new protocols. nftables is configured via the user-space utility *nft*, while legacy tools are configured via the utilities *iptables*, *ip6tables*, *arptables* and *ebtables* frameworks.

nftables utilizes the building blocks of the Netfilter infrastructure, such as the existing hooks into the networking stack, connection tracking system, userspace queueing component, and logging subsystem.

<https://en.wikipedia.org/wiki/Nftables>

nftables

/etc/nftables

To add config from file do

```
nft -f [filename].conf
```

To drop invalid traffic and allow out

```
ct state invalid counter drop comment "INVALID PACKET DROP"  
ct state { established, related } counter accept comment "RELATED CONNECTION ACCEPT"
```

Close bruteforce ssh

```
tcp dport ssh ct state new limit rate 15/minute accept
```

To include

```
/etc/sysconfig/nftables.conf
```

<https://en.wikipedia.org/wiki/Nftables>

<https://voxlink.ru/kb/linux/configuring-nftables-on-centos-8-for-ip-telephony-security/>

nftables

https://wiki.nftables.org/wiki-nftables/index.php/Main_Page

<https://en.wikipedia.org/wiki/Nftables>

quagga

```
sudo apt-get install quagga && sudo mkdir -p /var/log/quagga && sudo chown quagga:quagga /var/log/quagga
```

Create the configuration files:


```
sudo nano /etc/quagga/bgpd.conf  
sudo nano /etc/quagga/isisd.conf  
sudo nano /etc/quagga/ospf6d.conf  
sudo nano /etc/quagga/ospfd.conf  
sudo nano /etc/quagga/pimd.conf  
sudo nano /etc/quagga/ripd.conf  
sudo nano /etc/quagga/ripngd.conf  
sudo nano /etc/quagga/vtysh.conf  
sudo nano /etc/quagga/zebra.conf  
sudo /etc/init.d/quagga restart
```

<https://wiki.ubuntu.com/JonathanFerguson/Quagga>

<https://github.com/krobus00/Tubes-JARKOM/wiki/Install-Quagga-di-Ubuntu-20.04>

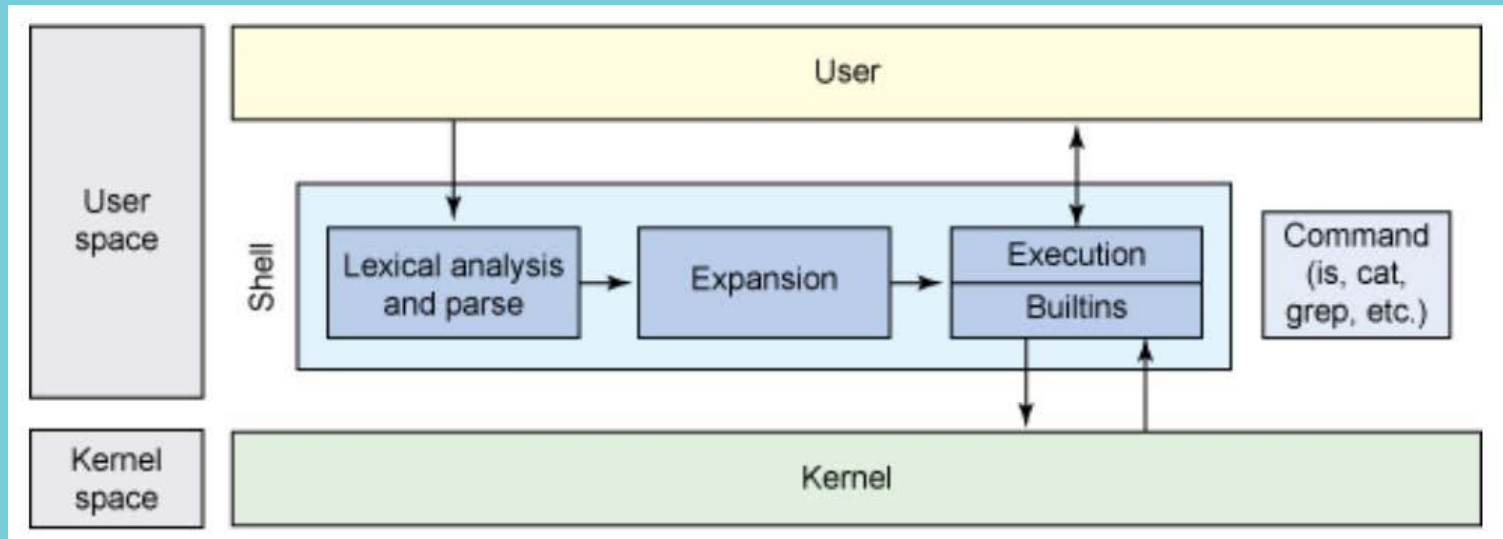
SHELL

Shell Evolution



Basic shell architecture

The fundamental architecture of a hypothetical shell is simple (as evidenced by Bourne's shell). As you can see below, the basic architecture looks similar to a pipeline, where input is analyzed and parsed, symbols are expanded (using a variety of methods such as brace, tilde, variable and parameter expansion and substitution, and file name generation), and finally commands are executed (using shell built-in commands, or external commands)



Bash. Scripting. Summary (1)

When **not** to use shell scripts:

- Resource-intensive tasks, especially where speed is a factor (sorting, hashing, recursion, etc.)
- Procedures involving heavy-duty math operations, especially floating point arithmetic, arbitrary precision calculations, or complex numbers (use C++ or FORTRAN instead)
- Cross-platform portability required (use C or Java instead)
- Complex applications, where structured programming is a necessity (type-checking of variables, function prototypes, etc.)
- Mission-critical applications upon which you are betting the future of the company
- Situations where security is important, where you need to guarantee the integrity of your system and protect against intrusion, cracking, and vandalism
- Project consists of subcomponents with interlocking dependencies
- Extensive file operations required (Bash is limited to serial file access)

Bash. Scripting. Summary (2)

When **not** to use shell scripts:

- Need native support for multi-dimensional arrays
- Need data structures, such as linked lists or trees
- Need to generate / manipulate graphics or GUIs
- Need direct access to system hardware or external peripherals
- Need port or socket I/O
- Need to use libraries or interface with legacy code
- Proprietary, closed-source applications (Shell scripts put the source code right out in the open for all the world to see.)

P.S. If any of the above applies, consider a more powerful scripting language -- perhaps Perl, Python, Ruby - or possibly a compiled language such as C, C++, or Java.

Even then, prototyping the application as a shell script might still be a useful development step.

Bash. Scripting. Summary (2)

Shells like **bash** have support for programming constructs that can be saved as scripts.

These scripts in turn then become more shell commands. Many Linux commands are scripts.

User profile scripts are run when a user logs on and init scripts are run when a daemon is stopped or started.

This means that system administrators also need basic knowledge of scripting to understand how their servers and their applications are started, updated, upgraded, patched, maintained, configured and removed, and also to understand how a user environment is built.

The goal of this module is to give enough information to be able to read and understand scripts. And to become a writer of simple scripts.

Bash. Scripting. Hello World

Just like in every programming course, we start with a simple `hello_world` script. The following script will output Hello World.

> echo Hello World

After creating this simple script in `vi` or with `echo`, you'll have to `chmod +x hello_world` to make it executable. And unless you add the scripts directory to your path, you'll have to type the path to the script for the shell to be able to find it.

```
[student@localhost ~]$ echo echo Hello World > hello_world
```

```
[student@localhost ~]$ chmod +x hello_world
```

```
[student@localhost ~]$ ./hello_world
```

```
Hello World
```

```
[student@localhost ~]$
```

Bash. Scripting. She-bang

Let's expand our example a little further by putting `#!/bin/bash` on the first line of the script. The `#!` is called a she-bang (sometimes called sha-bang), where the she-bang is the first two characters of the script.

```
#!/bin/bash
```

```
echo Hello World
```

You can never be sure which shell a user is running. A script that works flawlessly in bash might not work in ksh, csh, or dash. To instruct a shell to run your script in a certain shell, you can start your script with a she-bang followed by the shell it is supposed to run in. This script will run in a bash shell.

```
#!/bin/bash
```

```
echo -n hello
```

```
echo A bash subshell `echo -n hello`
```

Bash. Scripting. Comments. Variables

Let's expand our example a little further by adding comment lines.

```
#!/bin/bash  
#  
# Hello World Script  
#  
echo Hello World
```

Here is a simple example of a variable inside a script.

```
#!/bin/bash  
#  
# simple variable in script  
#  
var1=3  
echo var1 = $var1
```

Bash. Scripting. Variables. Sourcing a script

Scripts can contain variables, but since scripts are run in their own shell, the variables do not survive the end of the script.

```
[student@localhost ~]$ ./simple_variable_in_script
```

```
var1 = 3
```

```
[student@localhost ~]$ echo $var1
```

```
[student@localhost ~]$
```

But we can force a script to run in the same shell, this is called **sourcing** a script (2 ways).

```
[student@localhost ~]$ source ./simple_variable_in_script
```

```
var1 = 3
```

```
[student@localhost ~]$ echo $var1
```

```
3
```

```
[student@localhost ~]$
```

```
[student@localhost ~]$ . ./simple_variable_in_script
```

```
var1 = 3
```

```
[student@localhost ~]$ echo $var1
```

```
3
```

```
[student@localhost ~]$
```

Bash. Scripting. Conditions and loops. test[]

The **test** command can test whether something is true or false. Let's start by testing whether 10 is greater than 55.

```
$ test 10 -gt 55 ; echo $?
```

```
1
```

The test command returns **1** if the **test** fails. And as you see in the next screenshot, **test** returns 0 when a test succeeds.

```
$ test 56 -gt 55 ; echo $?
```

```
0
```

If you prefer true and false, then write the test like this.

```
$ test 56 -gt 55 && echo true || echo false
```

```
true
```

```
$ test 6 -gt 55 && echo true || echo false
```

```
false
```

The test command can also be written as square brackets, the screenshot below is identical to the one above.

```
$ [ 56 -gt 55 ] && echo true || echo false
```

```
true
```

```
$ [ 6 -gt 55 ] && echo true || echo false
```

```
false
```


Bash. Scripting. Conditions and loops. test[]

Below are some example tests. Take a look at *man test* to see more options for tests.

[-d foo] Does the directory foo exist ?

[-e bar] Does the file bar exist ?

['/etc' = \$PWD] Is the string /etc equal to the variable \$PWD ?

[\$1 != 'secret'] Is the first parameter different from secret ?

[55 -lt \$bar] Is 55 less than the value of \$bar ?

[\$foo -ge 1000] Is the value of \$foo greater or equal to 1000 ?

["abc" < \$bar] Does abc sort before the value of \$bar ?

[-f foo] Is foo a regular file ?

[-r bar] Is bar a readable file ?

[foo -nt bar] Is file foo newer than file bar ?

[-o nounset] Is the shell option nounset set ?

Tests can be combined with logical AND and OR.

*\$ [66 -gt 55 -a 66 -lt 500] && echo true || echo false
true*

*\$ [66 -gt 55 -a 660 -lt 500] && echo true || echo false
false*

*\$ [66 -gt 55 -o 660 -lt 500] && echo true || echo false
true*

Bash. Scripting. Conditions *If then else*

The *if then else* construction is about choice. If a certain condition is met, then execute something, else execute something else. The example below tests whether a file exists, and if the file exists then a proper message is echoed.

```
#!/bin/bash
```

```
if [ -f isit.txt ]
```

```
    then echo isit.txt exists!
```

```
else echo isit.txt not found!
```

```
fi
```

If we name the above script 'choice', then it executes like this.

```
$ ./choice
```

```
isit.txt not found!
```

```
$ touch isit.txt
```

```
$ ./choice isit.txt exists!
```

```
$
```

Bash. Scripting. Conditions *If then elif*

You can nest a new *if* inside an *else* with *elif*. This is a simple example.

```
#!/bin/bash
count=42
if [ $count -eq 42 ]
then
echo "42 is correct."
elif [ $count -gt 42 ]
then
echo "Too much."
else
echo "Not enough."
fi
```

Bash. Scripting. Loops. *for loop*

The example below shows the syntax of a classical *for loop* in bash:

```
for i in 1 2 4
```

```
do
```

```
echo $i
```

```
done
```

An example of a for loop combined with an embedded shell:

```
#!/bin/bash
```

```
for counter in `seq 1 20`
```

```
do
```

```
echo counting from 1 to 20, now at $counter
```

```
sleep 1
```

```
done
```

Bash. Scripting. Loops. *for loop*

The same example as above can be written without the embedded shell using the bash {from..to} shorthand.

```
#!/bin/bash
```

```
for counter in {1..20}
```

```
do
```

```
echo counting from 1 to 20, now at $counter
```

```
sleep 1
```

```
done
```

This for loop uses file globbing (from the shell expansion). Putting the instruction on the command line has identical functionality.

```
$ ls
```

```
count.ksh go.ksh
```

```
$ for file in *.ksh ; do cp $file $file.backup ; done
```

```
$ ls
```

```
count.ksh count.ksh.backup go.ksh go.ksh.backup
```

Bash. Scripting. Loops. *while* loop

Below a simple example of a ***while* loop**

Endless loops can be made with ***while true*** or ***while :***, where the ***colon*** is the equivalent of ***no operation*** in the ***bash*** shell.

Below a simple example of an ***until* loop**

```
i=100;
while [ $i -ge 0 ] ;
do
    echo Counting down, from 100 to 0, now at $i;
    let i--;
done
```

```
#!/bin/bash
# endless loop while :
do
    echo hello
    sleep 1
done
```

```
let i=100;
until [ $i -le 0 ] ;
do
    echo Counting down, from 100 to 1, now at $i;
    let i--;
done
```

Bash. Scripting. Loops. *for loop*

Write a script that counts the number of files ending in **.txt** in the current directory

```
#!/bin/bash
let i=0
for file in *.txt
do
    let i++
done
echo "There are $i files ending in .txt"
```

Wrap an **if** statement around the script so it is also correct when there are zero files ending in **.txt**

```
#!/bin/bash
ls *.txt > /dev/null 2>&1
if [ $? -ne 0 ]
then echo "Directory contains 0 *.txt files"
else
    let i=0
    for file in *.txt
    do
        let i++
    done
    echo " Directory contains $i *.txt files "
fi
```

QUESTIONS & ANSWERS

A world map with a light beige background and dark beige landmasses. The text "THANK YOU!" is centered over the Atlantic Ocean in a black, serif, all-caps font.

THANK YOU!